

AI System Audit Report

EU AI Act compliance assessment — Claude, Gemini, GPT-5, Llama 3.1 · 25 Jun 2026

24 Total findings	13 Gaps identified	13 High severity	24 Auditor reviewed
----------------------	-----------------------	---------------------	------------------------

Compliance heatmap

	claude	google_gemini	gpt5	meta_llama
RC-01	gap high	gap high	gap high	gap high
RC-02	partial medium	partial medium	partial medium	partial medium
RC-03	gap high	gap high	gap high	gap high
RC-04	partial medium	partial medium	gap high	partial medium
RC-05	gap high	gap high	gap high	gap high
RC-06	partial medium	partial medium	partial medium	partial medium

Color indicates severity: red = high, amber = medium, green = low/compliant.

Detailed findings

Claude

RC-01 — Risk management system		GAP	HIGH
Finding	The Responsible Scaling Policy is a policy commitment document, not a documented lifecycle risk management process as required by Article 9.		
Evidence	"focusing on core capabilities, safety, societal impacts, and the catastrophic risk assessments we committed to in our Responsible Scaling Policy"		
LLM suggestion	Develop and document a comprehensive risk management process covering the full lifecycle.		
Auditor rec.	Anthropic should produce a standalone lifecycle risk management document mapping risks at each stage: development, deployment, and post-deployment monitoring. The Responsible Scaling Policy alone is insufficient for Article 9 compliance.		
Reviewed	Yes		

RC-02 — Data governance

PARTIAL

MEDIUM

Finding	The document describes data sources and some governance practices but lacks comprehensive documentation.
Evidence	<i>"We employ several data cleaning and filtering methods, including deduplication and classification."</i>
LLM suggestion	Provide detailed documentation of data governance practices for training, validation, and test datasets.
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.
Reviewed	Yes

RC-03 — Technical documentation

GAP

HIGH

Finding	The document redirects to external research papers for comprehensive insights.
Evidence	<i>"For comprehensive insights into our training and evaluation methodologies, we invite you to explore our research papers."</i>
LLM suggestion	Include self-contained technical documentation within the document for regulatory assessment.
Auditor rec.	Model card must be self-contained for regulatory purposes. Anthropic should either consolidate key technical details inline or formally reference a publicly accessible annex that satisfies Article 11 requirements.
Reviewed	Yes

RC-04 — Transparency & instructions for use

PARTIAL

MEDIUM

Finding	The document provides some transparency but lacks detailed user instructions.
Evidence	<i>"We clearly document the ways in which our products may and may not be used, as well as the limitations and potential risks of using our products."</i>
LLM suggestion	Enhance transparency by providing detailed instructions for use and interpretation of outputs.
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.
Reviewed	Yes

RC-05 — Human oversight		GAP	HIGH
Finding	Human review is described as an internal QA mechanism for classifier accuracy, not an operator-level override capability as required by Article 14.		
Evidence	<i>"We use a combination of automated monitoring for potential harms and violations of our AUP, as well as human review to audit the accuracy of our classifiers."</i>		
LLM suggestion	Document detailed human oversight mechanisms, including intervention and override capabilities.		
Auditor rec.	Anthropic must document explicit operator-level override mechanisms separate from internal classifier review. This should include who can pause the system, under what conditions, and how.		
Reviewed	Yes		

RC-06 — Accuracy, robustness & cybersecurity		PARTIAL	MEDIUM
Finding	The document mentions evaluations for safety but lacks comprehensive details on maintaining performance under adversarial conditions.		
Evidence	<i>"We regularly evaluate our systems through interactive red teaming, as well as assessments against benchmarks for both product performance and potential safety risks."</i>		
LLM suggestion	Provide detailed documentation on maintaining system performance under normal and adversarial conditions.		
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.		
Reviewed	Yes		

Google Gemini

RC-01 — Risk management system		GAP	HIGH
Finding	No documented, continuous risk management process covering the full lifecycle is mentioned.		
Evidence	<i>"none found"</i>		
LLM suggestion	Implement and document a continuous risk management process covering the full lifecycle.		
Auditor rec.	Google DeepMind should publish a lifecycle risk management process covering pre-training, fine-tuning, deployment, and post-deployment monitoring phases.		
Reviewed	Yes		

RC-02 — Data governance

PARTIAL

MEDIUM

Finding	Training data processing is mentioned but lacks comprehensive governance documentation.
Evidence	<i>"Training videos were also filtered for various compliance, safety, and quality metrics and deduplicated semantically."</i>
LLM suggestion	Provide detailed documentation on data governance practices for training, validation, and test datasets.
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.
Reviewed	Yes

RC-03 — Technical documentation

GAP

HIGH

Finding	Document redirects to an external source rather than providing self-contained compliance documentation.
Evidence	<i>"See the Google DeepMind site for a comprehensive list of model cards."</i>
LLM suggestion	Include comprehensive technical documentation within the document itself.
Auditor rec.	Redirecting to an external model card list is insufficient under Article 11. Each model card must be self-contained or accompanied by a linked compliance annex.
Reviewed	Yes

RC-04 — Transparency & instructions for use

PARTIAL

MEDIUM

Finding	Some transparency is provided, but lacks detailed instructions for interpreting outputs.
Evidence	<i>"Model cards are intended to provide developers with essential, summarized information on models."</i>
LLM suggestion	Enhance transparency by providing detailed instructions for interpreting model outputs.
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.
Reviewed	Yes

RC-05 — Human oversight

GAP

HIGH

Finding	No mention of human oversight, intervention, or override capabilities.
Evidence	<i>"none found"</i>
LLM suggestion	Develop and document mechanisms for human oversight, intervention, and override.
Auditor rec.	No override or intervention mechanism documented anywhere in the card. Google DeepMind must document how operators can halt, pause, or override the system in deployment.
Reviewed	Yes

RC-06 — Accuracy, robustness & cybersecurity		PARTIAL	MEDIUM
Finding	Some safety measures are mentioned, but comprehensive performance maintenance under all conditions is not detailed.		
Evidence	<i>"A range of evaluations and red teaming activities were conducted to help improve the model safety."</i>		
LLM suggestion	Provide detailed documentation on how the system maintains performance under normal and adversarial conditions.		
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.		
Reviewed	Yes		

Gpt5

RC-01 — Risk management system		GAP	HIGH
Finding	The document does not provide a documented, continuous risk management process.		
Evidence	<i>"none found"</i>		
LLM suggestion	Develop and document a comprehensive risk management process covering the full lifecycle.		
Auditor rec.	Preparedness Framework references are promising but insufficient. OpenAI should produce a full Article 9-compliant risk register with lifecycle coverage and residual risk disclosures.		
Reviewed	Yes		

RC-02 — Data governance		PARTIAL	MEDIUM
Finding	The document mentions data filtering processes but lacks comprehensive governance practices.		
Evidence	<i>"Our data processing pipeline includes rigorous filtering to maintain data quality and mitigate potential risks."</i>		
LLM suggestion	Expand documentation to include detailed governance practices for training, validation, and test datasets.		
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.		
Reviewed	Yes		

RC-03 — Technical documentation

GAP

HIGH

Finding	The document lacks sufficient technical documentation for regulatory compliance assessment.
Evidence	"none found"
LLM suggestion	Provide detailed technical documentation that allows for regulatory compliance assessment.
Auditor rec.	Technical documentation gap is the most severe finding in this document. OpenAI must provide self-contained technical specs — architecture, training methodology, evaluation results — within the system card itself.
Reviewed	Yes

RC-04 — Transparency & instructions for use

GAP

HIGH

Finding	The document does not provide sufficient transparency for users to interpret outputs correctly.
Evidence	"none found"
LLM suggestion	Include clear instructions and transparency measures to help users interpret system outputs.
Auditor rec.	Usage policy URL reference does not satisfy Article 13. Inline instructions for interpreting model outputs and known failure modes must be present in the system card.
Reviewed	Yes

RC-05 — Human oversight

GAP

HIGH

Finding	The document does not mention mechanisms for human oversight, intervention, or override.
Evidence	"none found"
LLM suggestion	Implement and document mechanisms for effective human oversight and intervention.
Auditor rec.	No human oversight mechanism documented. OpenAI must specify operator controls, escalation paths, and override procedures as required by Article 14.
Reviewed	Yes

RC-06 — Accuracy, robustness & cybersecurity		PARTIAL	MEDIUM
Finding	The document discusses robustness evaluations but lacks comprehensive coverage of accuracy and cybersecurity.		
Evidence	<i>"We subjected the model to our full suite of predeployment safety evaluations and our Preparedness Framework."</i>		
LLM suggestion	Enhance documentation to cover accuracy and cybersecurity measures comprehensively.		
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.		
Reviewed	Yes		

Meta Llama

RC-01 — Risk management system		GAP	HIGH
Finding	The document does not mention a risk management process covering the full lifecycle.		
Evidence	<i>"none found"</i>		
LLM suggestion	Implement and document a comprehensive risk management process covering the full lifecycle of the AI system.		
Auditor rec.	Meta provides no lifecycle risk management documentation. Given Llama's open-source distribution model, this is especially critical — downstream deployers have no risk baseline to build from.		
Reviewed	Yes		

RC-02 — Data governance		PARTIAL	MEDIUM
Finding	The document mentions data sources and freshness but lacks detailed governance practices.		
Evidence	<i>"Llama 3.1 was pretrained on ~15 trillion tokens of data from publicly available sources."</i>		
LLM suggestion	Provide detailed documentation on data governance practices, including training, validation, and test datasets.		
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.		
Reviewed	Yes		

RC-03 — Technical documentation

GAP

HIGH

Finding	The document redirects to external sources for technical details.
Evidence	<i>"For more technical information about generation parameters and recipes for how to use Llama 3.1 in applications, please go here."</i>
LLM suggestion	Include comprehensive technical documentation within the document itself to ensure self-contained compliance.
Auditor rec.	Redirecting to GitHub for technical details does not satisfy Article 11. Meta must publish a self-contained compliance document, particularly given the open-source deployment context.
Reviewed	Yes

RC-04 — Transparency & instructions for use

PARTIAL

MEDIUM

Finding	The document provides some intended use cases but lacks detailed instructions for interpreting outputs.
Evidence	<i>"Llama 3.1 is intended for commercial and research use in multiple languages."</i>
LLM suggestion	Enhance transparency by providing detailed instructions for interpreting outputs and using the system.
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.
Reviewed	Yes

RC-05 — Human oversight

GAP

HIGH

Finding	The document does not mention mechanisms for human oversight, intervention, or override.
Evidence	<i>"Developers are expected to deploy system safeguards when building agentic systems."</i>
LLM suggestion	Develop and document mechanisms for effective human oversight, intervention, and override of the AI system.
Auditor rec.	Delegating oversight entirely to downstream developers without baseline requirements creates an accountability gap. Meta should define minimum human oversight requirements for all Llama deployments.
Reviewed	Yes

RC-06 — Accuracy, robustness & cybersecurity		PARTIAL	MEDIUM
Finding	The document mentions performance benchmarks but lacks details on robustness and cybersecurity measures.		
Evidence	<i>"In this section, we report the results for Llama 3.1 models on standard automatic benchmarks."</i>		
LLM suggestion	Include detailed documentation on measures taken to ensure robustness and cybersecurity of the AI system.		
Auditor rec.	Partial compliance noted. Provider should expand documentation to fully satisfy the relevant EU AI Act article. Re-assess at next audit cycle.		
Reviewed	Yes		

Systemic recommendations

RC-05 — gap in 4/4 documents

Human oversight is the most systemic gap — absent in 3 of 4 documents. All providers should publish operator-level override and intervention procedures as a minimum baseline for EU AI Act Article 14 compliance.

RC-03 — gap in 4/4 documents

Technical documentation is redirected to external sources across all 4 documents. Regulators cannot rely on documents that defer compliance evidence elsewhere. Providers should produce self-contained compliance annexes.

RC-01 — gap in 4/4 documents

Risk management lifecycle documentation is absent in 3 documents. Providers must distinguish between policy commitments and operational risk management processes — the former does not satisfy Article 9.

Limitations

This assessment is based solely on publicly available model cards and system cards. Findings reflect documentation quality, not necessarily the underlying system's actual compliance posture. LLM-generated findings were reviewed and overridden where necessary by a human auditor. A full EU AI Act compliance assessment would require access to internal technical documentation, audit logs, and direct engagement with the provider.